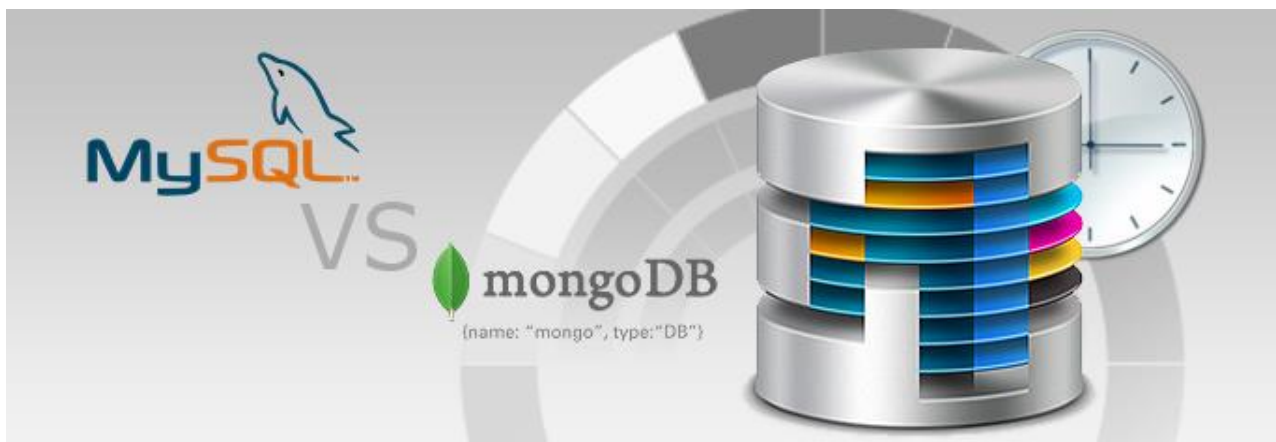




GUTENBERG PROJECT REPORT

Made by Murched Kayed, Hallur vid Neyst & Zaeem Shafiq



MAY 30, 2019

COPENHAGEN BUSINESS ACADEMY

Table of Contents

Introduction	2
Databases used	2
Programming language used	2
Data modeling in the database	2
MySQL	2
Indexes	3
MongoDB	4
Data modeling in the application	5
Data import	6
Behavior of query test set	9
Query runtime influenced by the database vs. the application frontend	9
Conclusion	10
Which database should be used in such a project for production?	10

Introduction

This report is about the solution we made for the Gutenberg project. We will, among other things, write about which databases and technologies we used to solve it. Furthermore, the report describes how we imported and structured the data in the databases.

Databases used

We have chosen to make use of two different kinds of databases: one relational database and one NoSQL database. We chose MySQL as the relational database and MongoDB as the NoSQL database.

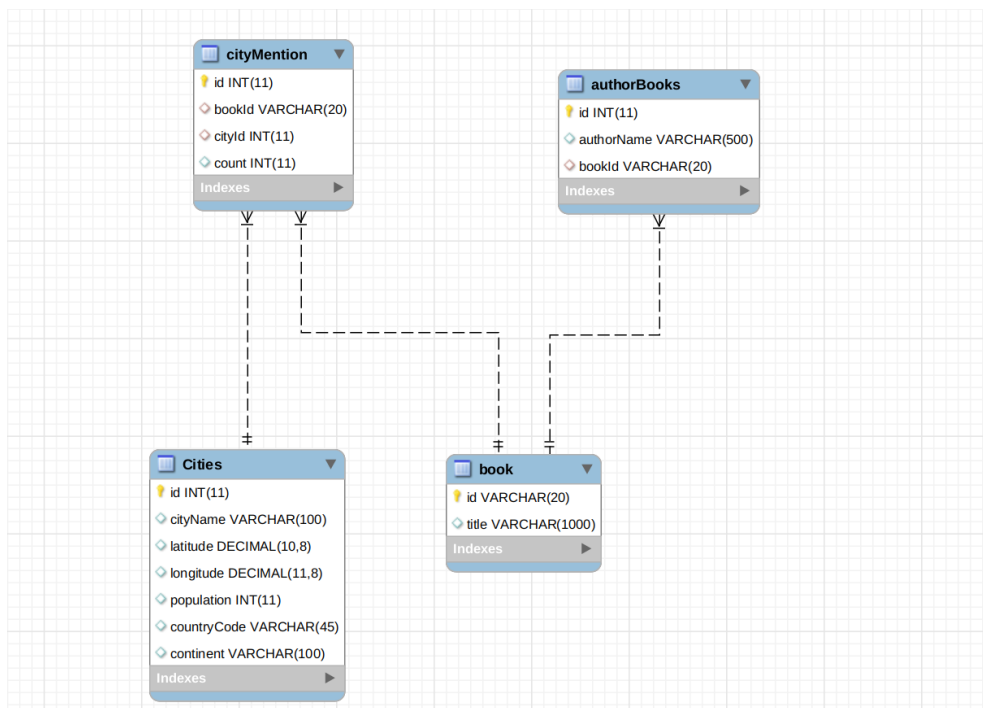
Programming language used

We agreed as a group that we will use Java to solve this project, since we all three had good skills in Java. To make it easier to handle dependencies, we have made our project using Maven. We have chosen to focus on the database part of the project and therefore just used a simple Command-line interface as our frontend.

Data modeling in the database

MySQL

We have chosen to structure our tables in MySQL as follows:

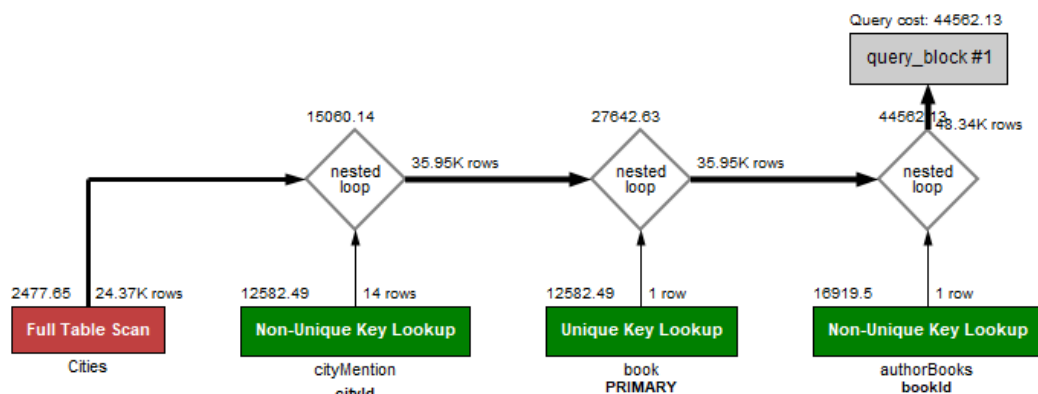


Indexes

To optimize performance, we have chosen to use indexes on some of our columns in the database. This has proved to be extremely useful and given us a much faster performance on our queries.

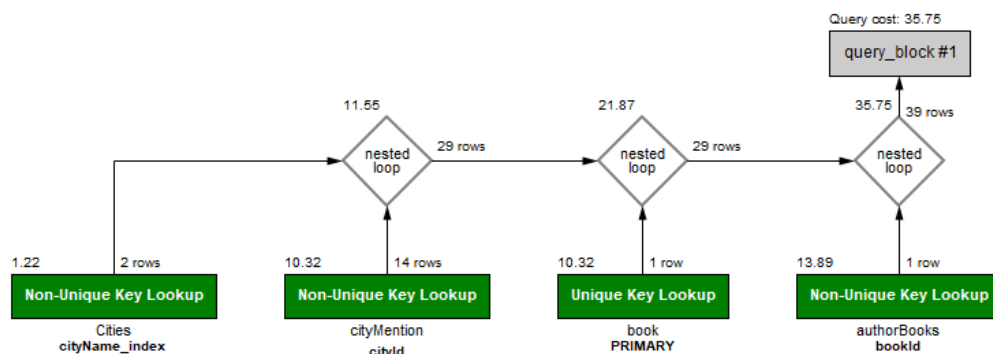
Below are a number of examples of the graphical execution plan where we make use of the same query both before and after we have used indexes.

Before creating the index *cityName_index*:

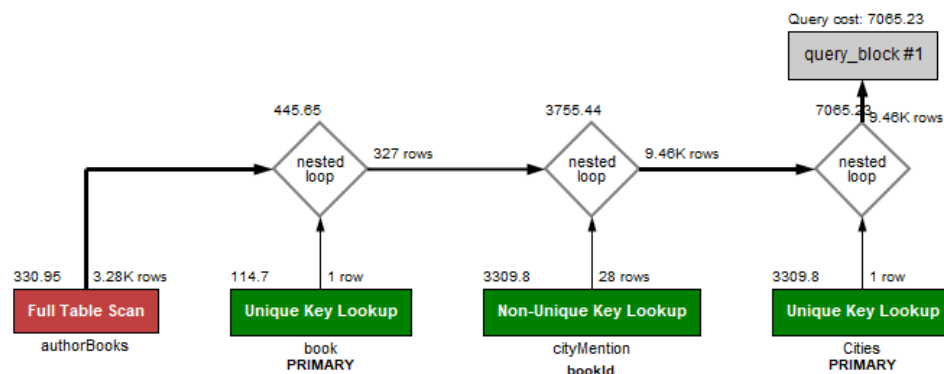


After creating the index:

```
create index cityName_index on Cities(cityName);
```

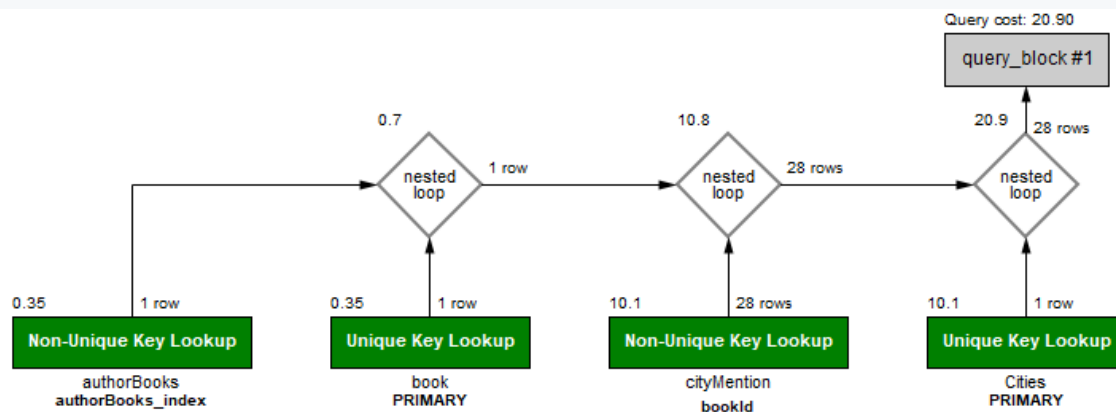


Before creating the index *authorName_index*:



After creating the index:

```
create index authorBooks_index on authorBooks(authorName);
```



As it can be seen above, our performance has been drastically improved after we implemented indexes in our MySQL database. We avoid making a Full Table Scan, which turns out to save us a lot of time. E.g. in one of the examples above the query cost goes from 44562.13 to only 35.75. This is a great improvement.

MongoDB

Since MongoDB is a NoSQL database and hence doesn't contain relations, we have chosen a completely different structure for our MongoDB database. We have chosen to simply store all the data in one collection. Our collection is structured as follows:

```

1 {
2   "_id" : ObjectId("5ceeb09a5b2567bcd4cf1afa"),
3   "authorName" : "Kevorkian, Hagop K.",
4   "books" : [
5     {
6       "title" : "The Arts of Persia & Other Countries of Islam",
7       "cities" : [
8         {
9           "cityName" : "Armenia",
10          "latitude" : "4.53389",
11          "longitude" : "-75.68111",
12          "count" : "1"
13        },
14        {
15          "cityName" : "Asia",
16          "latitude" : "9.5506",
17          "longitude" : "122.5164",
18          "count" : "1"
19        },
20        {
21          "cityName" : "Paris",
22          "latitude" : "33.66094",
23          "longitude" : "-95.55551",
24          "count" : "2"
25        }
26      ]
27    }
28  ]
29 }

```

As seen in the picture above, we only have one collection in our database. This collection follows the following structure: In the object is an author. An author contains an ID, a name, and a list of the books he/she has written. The list of books then contains a title and a list of the cities mentioned in that book. The list of cities contains different attributes about a city: the city's name, location and how many times the city is mentioned in the book (count). This simple structure has been followed throughout the whole collection.

Data modeling in the application


To have a better overview of our project, we have structured our classes into different packages as follows:

- **entity**
 - This package simply contains all our entity classes with corresponding attributes.
- **entitymanager**

- This package contains classes with methods that help to read books/cities from the txt/RDF files. It also contains methods that generate csv files with a header and all necessary attributes. These csv files can later be used to import data into our databases.
- ***files***
 - This package contains all the files we have used to create our databases and to insert data into our databases.
- ***main***
 - This package just contains our main method that runs our program.
- ***mongodb***
 - This package contains all communication between our program and the MongoDB database. That is, it contains both a connector that connects to the database and a number of methods that send the various queries to the MongoDB database.
- ***mysql***
 - This package contains all communication between our program and the MySQL database. That is, it contains both a connector that connects to the database and a number of methods that send the various queries to the MySQL database.
- ***threads***
 - This package contains eight threads classes. These have been made to optimize the process when looking for city names in our books (txt files). Since we worked with a lot of books/cities, the program was very slow as it had to read the files. We therefore chose to run eight threads at the same time, which proved to improve performance a lot.

Data import

Below is a picture of how we have imported the data in our program.



As it can be seen on the diagram above, we created a Java program to import the data into the databases. Our program works as follows:

1. Our program starts by reading the books/cities from different files. This is done as follows:
 - The cities and their geolocations are read from the csv file found here: <http://download.geonames.org/export/dump/cities15000.zip>. For this we used a Java library called opencsv (<http://opencsv.sourceforge.net/>).
 - The author names and book titles are read from the XML/RDF files from the offline catalogue available here: <https://www.gutenberg.org/wiki/Gutenberg:Feeds>. For this we used a Java library called Apache Jena (<https://jena.apache.org/documentation/io/rdf-input.html>).

- The cities mentioned in a book are found in the txt files we downloaded from our droplet on Digitalocean which downloaded them from the Gutenberg webpage. To recognize the cities mentioned in a book, we used the Stanford Named Entity Recognizer (NER) in Java (<https://nlp.stanford.edu/software/CRF-NER.html>).

2. Then our program generates a CSV/JSON file that can be used to import the data into the databases. The files only contain the attributes we agreed were needed to solve the queries. This is done as follows:

- We created csv files (separated by tab characters) that can be easily imported in MySQL. An example of how a part of a csv file looks like is (the first row is the header):

cityId	cityName	latitude	longitude	population	countryCode	continent
3040051	les Escaldes	42.50729	1.53414 15853	AD	Europe/Andorra	
3041563	Andorra la Vella	42.50779	1.52109	20430	AD	Europe/Andorra
290594	Umm al Qaywayn	25.56473	55.55517	44411	AE	Asia/Dubai
291074	Ras al-Khaimah	25.78953	55.9432 115949	AE	Asia/Dubai	
291696	Khawr Fakkān	25.33132	56.34199	33575	AE	Asia/Dubai
292223	Dubai	25.0657 55.17128	1137347	AE	Asia/Dubai	
292231	Dibba Al-Fujairah	25.59246	56.26176	30000	AE	Asia/Dubai
292239	Dibba Al-Hisn	25.61955	56.27291	26395	AE	Asia/Dubai
292672	Sharjah	25.33737	55.41206	1324473	AE	Asia/Dubai
292688	Ar Ruways	24.11028	52.73056	16000	AE	Asia/Dubai
292878	Al Fujayrah	25.11641	56.34141	62415	AE	Asia/Dubai
292913	Al Ain	24.19167	55.76056	408733	AE	Asia/Dubai
292932	Ajman	25.41111	55.43504	226172	AE	Asia/Dubai
292953	Adh Dhayd	25.28812	55.88157	24716	AE	Asia/Dubai
292968	Abu Dhabi	24.46667	54.36667	603492	AE	Asia/Dubai
12042053	Musaffah	24.35893	54.48267	243341	AE	Asia/Dubai
1120985	Zaranj	30.95962	61.86037	49851	AF	Asia/Kabul
1123004	Taloqan	36.73605	69.53451	64256	AF	Asia/Kabul
1125155	Shīngand	33.30294	62.1474 29264	AF	Asia/Kabul	
1125444	Shibirghān	36.66757	65.7529 55641	AF	Asia/Kabul	
1125896	Shahrak	34.10737	64.3052 15967	AF	Asia/Kabul	
1127110	Sar-e Pul	36.21544	65.93249	52121	AF	Asia/Kabul
1127628	Sang-e Chārak	35.84972	66.43694	15377	AF	Asia/Kabul
1127768	Aībak	36.26468	68.01551	47823	AF	Asia/Kabul
1128265	Rustāq	37.12604	69.83045	25636	AF	Asia/Kabul
1129516	Qarqīn	37.41853	66.04358	15018	AF	Asia/Kabul
1129648	Qarāwul	37.21959	68.7802 24544	AF	Asia/Kabul	
1130490	Pul-e Khumri	35.94458	68.71512	56369	AF	Asia/Kabul
1131316	Paghmān	34.58787	68.95091	49157	AF	Asia/Kabul
1132495	Nahrīn	36.0649 69.13343	22363	AF	Asia/Kabul	

- We created a JSON file that can be easily imported in MongoDB. An example of how a part of the JSON file looks like is:

```

(authorName: "Diffin, Charles Willard",
books: [
  {title: "The Hammer of Thor", cities: [
    {cityName: "Chicago",latitude : "41.85003", longitude: "-87.65005", count: "1"
    },
    {cityName: "Hudson",latitude : "41.24006", longitude: "-81.44067", count: "1"
    },
    {cityName: "Cleveland",latitude : "41.4995", longitude: "-81.69541", count: "1"
    },
    {cityName: "Washington",latitude : "37.13054", longitude: "-113.50825", count: "7"
    },
    {cityName: "Boston",latitude : "42.35843", longitude: "-71.05977", count: "3"
    }
  ]
},
],
),
(authorName: "Emerson, Charles Wesley",
books: [
  {title: "Evolution of Expression, Volume 2-RevisedA Compilation of Selections Illustrating the Four Stages of Development in Art As Applied to Oratory: Twenty-Eighth Edition", cities: [
    {cityName: "Rome",latitude : "43.21285", longitude: "-75.45573", count: "6"
    },
    {cityName: "Providence",latitude : "41.82399", longitude: "-71.41283", count: "1"
    },
    {cityName: "London",latitude : "51.50853", longitude: "-0.12574", count: "1"
    },
    {cityName: "Concord",latitude : "37.97798", longitude: "-122.03107", count: "3"
    },
    {cityName: "Winchester",latitude : "36.12997", longitude: "-115.11889", count: "1"
    },
    {cityName: "Lexington",latitude : "42.44732", longitude: "-71.2245", count: "4"
    },
    {cityName: "Washington",latitude : "37.13054", longitude: "-113.50825", count: "1"
    },
    {cityName: "Paris",latitude : "33.66094", longitude: "-95.55551", count: "1"
    },
    {cityName: "Rialto",latitude : "34.1064", longitude: "-117.37032", count: "1"
    },
    {cityName: "Boston",latitude : "42.35843", longitude: "-71.05977", count: "1"
    }
  ]
},
],
),

```

3. Then, via our program, we can easily insert the data into our databases by using the generated files. This is done as follows:

- In MySQL the data can be imported using the following query:

```

LOAD DATA INFILE '/home/bookTable.csv'
INTO TABLE book
FIELDS TERMINATED BY '\t'
LINES TERMINATED BY '\n'
IGNORE 1 ROWS;

```

- In MongoDB the data can be imported using the following command:

```

mongoimport --db your db Name --collection authors --file authorsJson.json --jsonArray

```

Behavior of query test set

Query runtime influenced by the database vs. the application frontend

It turns out that there is a difference in how long it takes to execute a query if you choose to run it directly in MySQL Workbench instead of running it through the program. This is because, among other things, Java must first connect to the database using a JDBC driver, whereas MySQL Workbench already has a connection to the database and thus can execute the query faster. The same goes for our MongoDB queries. For example, Robo 3T executes the queries faster than Java. In addition, we also found out that MySQL Workbench uses a

default limit of 1000 rows, whereas Java doesn't have any limit. This of course also made our queries run much faster in MySQL Workbench.

Conclusion

All in all, the whole development process has been extremely instructive for all of us. We experienced a lot of exciting challenges along the way and it has also given us the opportunity to try a lot of things on our own. We have become much wiser on the use of indexes and seen how much importance they can have for the query performance when working with a lot of data. In addition, we have strengthened our competences in reading from and writing to files of various formats.

Which database should be used in such a project for production?

Since we did not use the same tables in both our databases, we actually had a good performance in both databases. If we had used the same structure in MongoDB, which we used in MySQL, our queries in MongoDB would have become a lot slower. This is due to the fact that we would have had four tables in MongoDB, which means that we would have to use joins even though MongoDB was not originally designed for joining collections. There is no doubt that both databases have their pros and cons, and which one is the best, greatly depends on what you want to store in the database.

If we were to recommend a database for a project like Gutenberg, we would recommend MongoDB. This is partly due to the fact that we could make much simpler queries in MongoDB, as we did not need to join with other collections. In addition, this project contained a lot of work with cities and their locations and here MongoDB also has the great advantage that you can make use of the Geospatial Index to store city locations. Furthermore, JSON is used in MongoDB, which is also a great advantage.